

# Letter Image Recognition

**Source:** David J. Slate(1991) | UCI ML repository



Group 17 : Jacob & Mrigank

# Data description

- Character images based on **20** different fonts
- Each letter was randomly **distorted** to produce a file of **20,000** unique stimuli
- Statistical moments and edge counts, **scaled** to fit into a range of integer values from 0 through 15
- Attributes: 16 | Target: 26 classes
- Uniform **class distribution** among 26 letters.
- picture borrowed from: Frey, Peter W. & Slate, David J., 1991, Letter Recognition Using Holland-Style Adaptive Classifiers



# Models / Algorithms

- Logistic Regression using sigmoid activations
- Neural Networks with a hidden layer of ReLU activations and SoftMax output.
- K-means / cut-off clustering
- Ensemble: Logistic with clustering

# Logistic Regression

(Multinomial regression)

Prediction = maximum likelihood




Problems:

1. Bad performance before normalization
2. Sigmoid does not play well with gradient descent (saturation)
3. Sensitive to small variations

---

# Gradient Descent with Momentum

(Karpathy, Andrej)



How to avoid finding local minima in the error function?

Nesterov momentum.

—

# Output Encoding

## → Logistic Regression

Actual targets: 'A', 'B', ..., 'Z'

Targets encoded to 26 digits initialized to 0s.

Ex. 'T' → `(array([ 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0.,  
0., 0., 0., 0., 0., 0., 1., 0., 0., 0., 0., 0.]),`

## → Neural Networks

Each of the 26 letters are assigned a number between 0-25, representing 26 different classes.

Ex. 'T' → 19

# Multilayer Perceptron


→ MLP Classifier with 1 hidden layer containing **ReLU** activations and the **SoftMax** output layer

→ **ReLU pros**: [Krizhevsky et al., 2012]

1. Does not saturate (in +ve regions)
2. Very computationally efficient
3. Converges much faster than sigmoid/tanh

→ **SoftMax** (generalization of sigmoid)

Squashes a k-dimensional vector to a k-D vector of real values(0,1) that add **up to 1**.



Left: ReLU activations | Right: the significant improvement in convergence with the ReLU unit compared to the tanh unit ([reference](#))

$$p_k = \frac{e^{f_k}}{\sum_j e^{f_j}}$$

$$L_i = -\log(p_{y_i})$$

# Network Learning

→ Tried training networks with different activations

## Hyperparameters


```
clf1.get_params()
```

```
{'activation': 'relu',  
 'alpha': 1e-05,  
 'batch_size': 'auto',  
 'beta_1': 0.9,  
 'beta_2': 0.999,  
 'early_stopping': False,  
 'epsilon': 1e-08,  
 'hidden_layer_sizes': (100,),  
 'learning_rate': 'constant',  
 'learning_rate_init': 0.001,  
 'max_iter': 200,  
 'momentum': 0.9,  
 'nesterovs_momentum': True,  
 'power_t': 0.5,  
 'random_state': 1,  
 'shuffle': True,  
 'solver': 'lbfgs',  
 'tol': 0.0001,  
 'validation_fraction': 0.1,  
 'verbose': False,  
 'warm_start': False}
```



# Network Learning

→ Tried training networks with different activations



## Hyperparameters

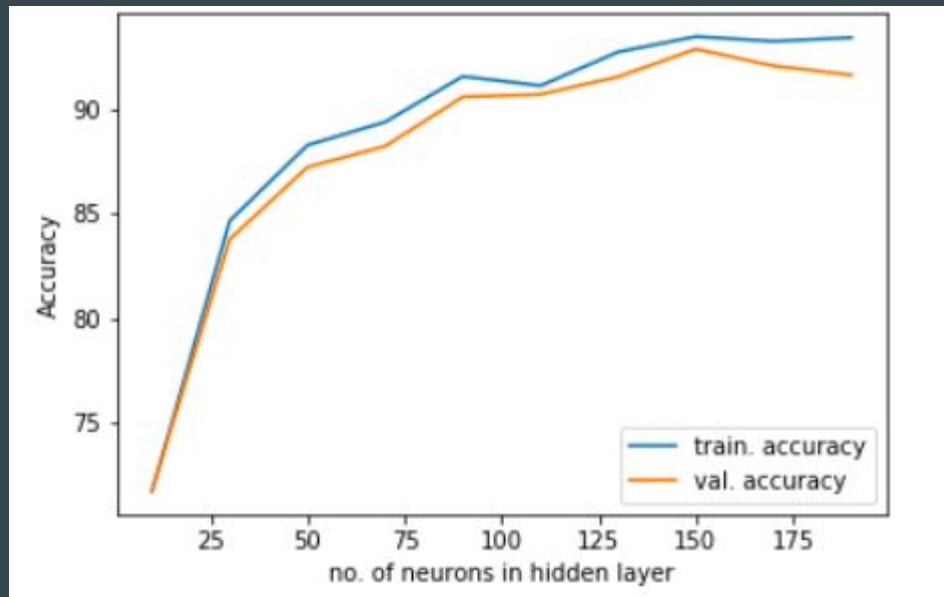
```
clf1.get_params()
```

```
{'activation': 'relu',  
 'alpha': 1e-05,  
 'batch_size': 'auto',  
 'beta_1': 0.9,  
 'beta_2': 0.999,  
 'early_stopping': False,  
 'epsilon': 1e-08,  
 'hidden_layer_sizes': (100,),  
 'learning_rate': 'constant',  
 'learning_rate_init': 0.001,  
 'max_iter': 200,  
 'momentum': 0.9,  
 'nesterovs_momentum': True,  
 'power_t': 0.5,  
 'random_state': 1,  
 'shuffle': True,  
 'solver': 'lbfgs',  
 'tol': 0.0001,  
 'validation_fraction': 0.1,  
 'verbose': False,  
 'warm_start': False}
```

# Prediction vs hyperparameters

→ Accuracy is on increasing trend with increase in number of hidden neurons

→ The white space between blue and yellow curve shows very **little over-fitting**.



# Clustering

K means

# clusters = 260

$\text{between\_class\_SS} / \text{total\_SS} = 0.857$

Average clusters per letter = 10


---

# Clustering

## Cutoff Clustering

Each point is clustered with its “neighbors”: points within a certain cutoff distance


Based on cluster density, not size



# Clustering

## Cutoff Clustering Results

Cutoff = 3.5



# Two-layer Ensemble

Clustering + Logistic Regression

Predictors = k-means label +  
cutoff-clustering label

# clusters in k-means = 26

Accuracy = 37%

---

# Performance

## Metrics

Accuracy

Maximal confusion rate

Average confusion rate

---

# Confusion Rate

	G	C		G	C
G	8	2	G	.8	.2
C	4	6	C	.4	.6

$$\begin{aligned}\text{Confusion rate} \\ &= 0.8 * 0.6 - 0.4 * 0.2 = 0.4\end{aligned}$$

$$= P(G | G) * P(C | C) - P(G | C) * P(C | G)$$

---



# Performance

Logistic Regression

Accuracy: 73%

MCR: 0.65 (confusing O and H)

ACR: 0.97

---

# Performance

Multilayer Perceptron

Accuracy: 95%

MCR: 0.88 (confusing G and C)

ACR: 0.9875

---

# Performance

Ensemble Model

Accuracy: 37%

MCR: 0.007 (confusing S and Z)

ACR: 0.875

---

# Conclusion

Letter recognition can be hard,  
especially for noisy data

Good confusion rate shows low  
systematic misprediction

---

# Questions?

Models:

1. Logistic Regression
2. Multilayer Perceptron
3. K-means clustering
4. Cutoff clustering

Metrics:

1. Accuracy
  2. Maximal Confusion Rate
  3. Average Confusion Rate
-